

# Computer Vision

Spring 2006 15-385,-685

Instructor: S. Narasimhan

Wean 5403

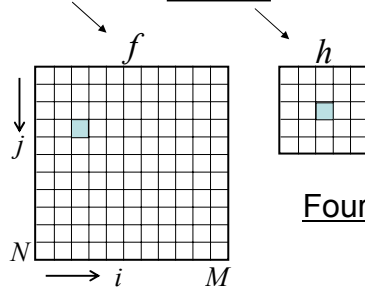
T-R 3:00pm – 4:20pm

Image Processing and Filtering  
(continued)

Lecture #6

# Images are Discrete and Finite

$$f(x, y) \rightarrow h(x, y) \rightarrow g(x, y)$$



## Convolution

$$g(i, j) = \sum_{m=1}^M \sum_{n=1}^N f(m, n) h(i-m, j-n)$$

## Fourier Transform

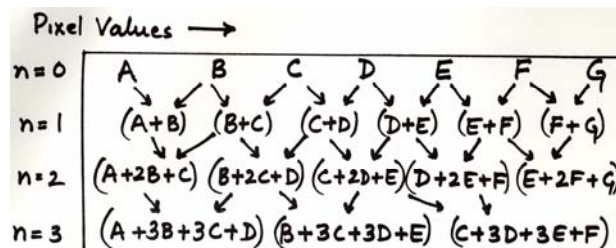
$$F(u, v) = \sum_{m=1}^M \sum_{n=1}^N f(m, n) e^{-i2\pi \left( \frac{mu}{M} + \frac{nv}{N} \right)}$$

## Inverse Fourier Transform

$$f(k, l) = \frac{1}{MN} \sum_{u=1}^M \sum_{v=1}^N F(u, v) e^{i2\pi \left( \frac{ku}{M} + \frac{lv}{N} \right)}$$

# Averaging

Let's think about averaging pixel values



For  $n=2$ , convolve pixel values with  $\begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$

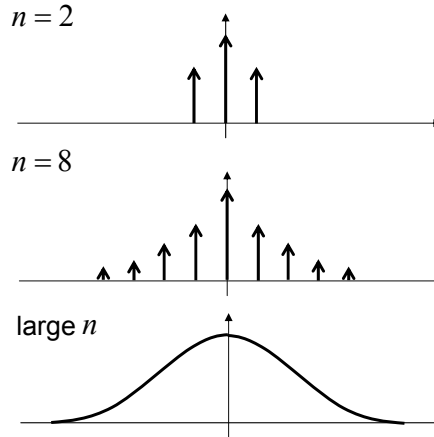
2D images:

(a) use  $\begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$  then  $\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix}$  or (b) use  $\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$

Which is faster? (a)  $O(2(n+1))$  (b)  $O((n+1)^2)$

# Averaging

The convolution kernel

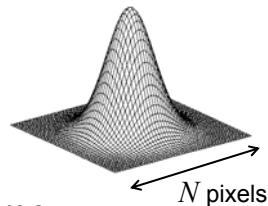


Repeated averaging  $\approx$  Gaussian smoothing

# Gaussian Smoothing

Gaussian kernel

$$h(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{1}{2}\left(\frac{i^2+j^2}{\sigma^2}\right)}$$



Filter size  $N \propto \sigma$  ...can be very large  
(truncate, if necessary)

$$g(i, j) = \frac{1}{2\pi\sigma^2} \sum_{m=1} \sum_{n=1} e^{-\frac{1}{2}\left(\frac{m^2+n^2}{\sigma^2}\right)} f(i-m, j-n)$$

2D Gaussian is separable!

$$g(i, j) = \frac{1}{2\pi\sigma^2} \sum_{m=1} e^{-\frac{1}{2}\frac{m^2}{\sigma^2}} \sum_{n=1} e^{-\frac{1}{2}\frac{n^2}{\sigma^2}} f(i-m, j-n)$$

Use two 1D Gaussian filters

## Gaussian Smoothing

- A Gaussian kernel gives less weight to pixels further from the center of the window

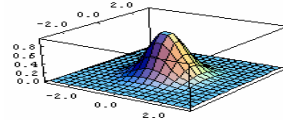
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$F[x, y]$

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} H[u, v]$$

- This kernel is an approximation of a Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$

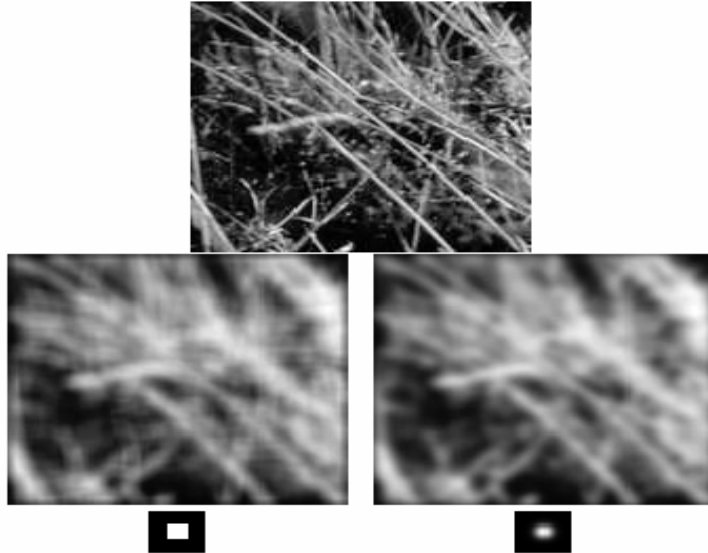


## Gaussian Smoothing



## Mean vs. Gaussian filtering

---



## Gaussian Smoothing

---



by Charles Allen Gillbert



by Harmon & Julesz

## Gaussian Smoothing

---



[http://www.michaelbach.de/ot/cog\\_blureffects/index.html](http://www.michaelbach.de/ot/cog_blureffects/index.html)

## Border Problem

---



## Border Problem

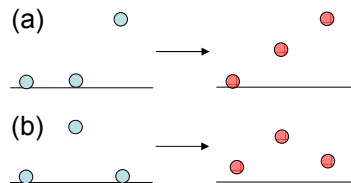
- Ignore
  - Output image will be smaller than original
- Pad with constant values
  - Can introduce substantial 1<sup>st</sup> order derivative values
- Pad with reflection
  - Can introduce substantial 2<sup>nd</sup> order derivative values

## Median Filter

- Smoothing is averaging

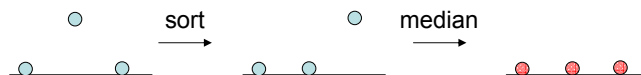
(a) Blurs edges

(b) Sensitive to outliers

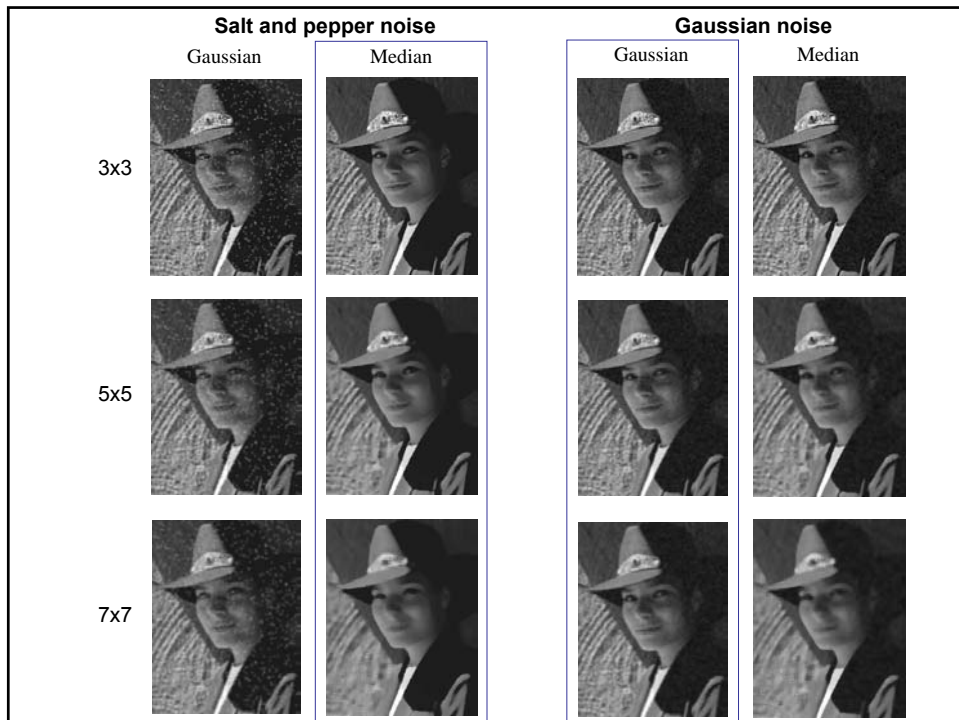


- Median filtering

- Sort  $N^2 - 1$  values around the pixel
- Select middle value (median)



- Non-linear (Cannot be implemented with convolution)



## Correlation



template

How do we locate the template in the image?

- Minimize

$$\begin{aligned}
 E(i, j) &= \sum_m \sum_n [f(m, n) - t(m - i, n - j)]^2 \\
 &= \sum_m \sum_n [f^2(m, n) + t^2(m - i, n - j) - 2f(m, n)t(m - i, n - j)]^2
 \end{aligned}$$

- Maximize

$$R_{tf}(i, j) = \sum_m \sum_n t(m - i, n - j)f(m, n)$$

Cross-correlation



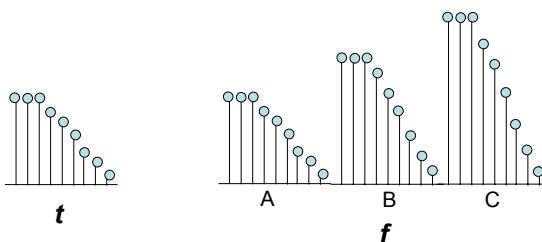
## Cross-correlation

$$R_{tf}(i, j) = \sum_m \sum_n t(m-i, n-j) f(m, n) \quad R_{tf} = t \otimes f$$

Note:  $t \otimes f \neq f \otimes t$

$$R_{ff} = f \otimes f \quad \text{Auto-correlation}$$

Problem:



$$R_{tf}(C) > R_{tf}(B) > R_{tf}(A) \quad \text{We need } R_{tf}(A) \text{ to be the maximum!}$$

## Normalized Correlation

- Account for energy differences

$$N_{tf}(i, j) = \frac{\sum_m \sum_n t(m-i, n-j) f(m, n)}{\left[ \sum_m \sum_n t^2(m-i, n-i) \right]^{1/2} \left[ \sum_m \sum_n f^2(m, n) \right]^{1/2}}$$

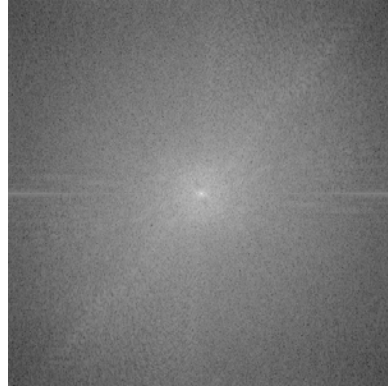


## Image Processing in the Fourier Domain

---



Magnitude of the FT



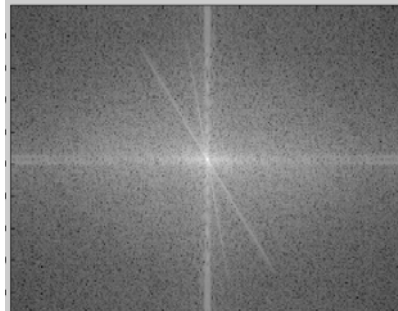
Does not look anything like what we have seen

## Image Processing in the Fourier Domain

---

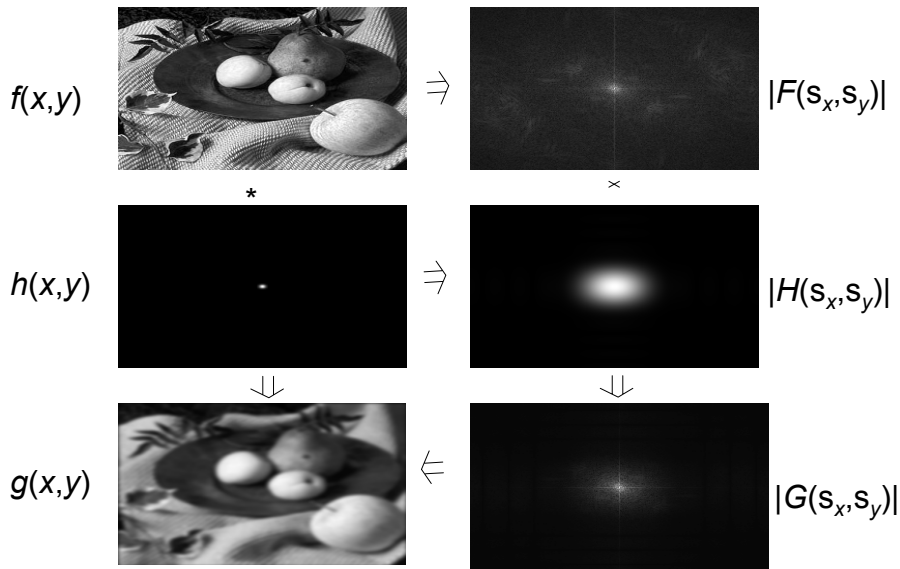


Magnitude of the FT

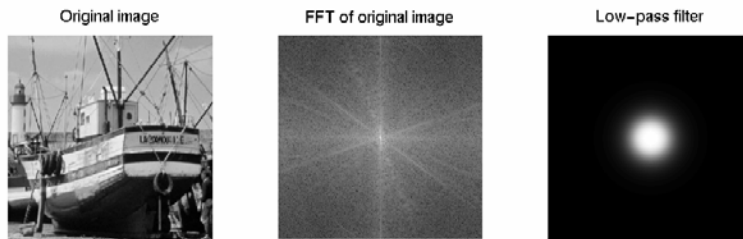


Does not look anything like what we have seen

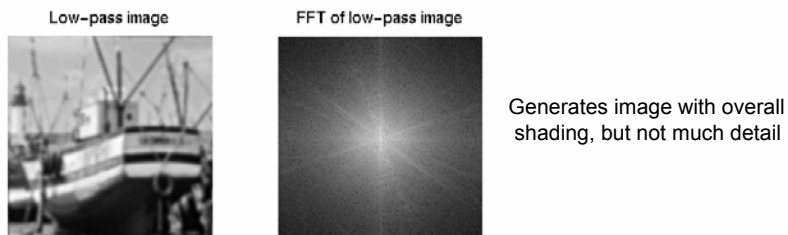
## Convolution is Multiplication in Fourier Domain



## Low-pass Filtering



Let the low frequencies pass and eliminating the high frequencies.

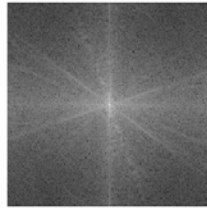


# High-pass Filtering

Original image



FFT of original image

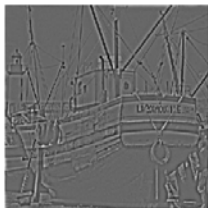


High-pass filter

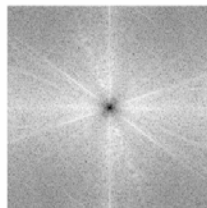


Lets through the high frequencies (the detail), but eliminates the low frequencies (the overall shape). It acts like an edge enhancer.

High-pass image



FFT of high-pass image

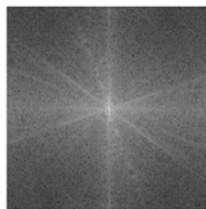


# Boosting High Frequencies

Original image



FFT of original image



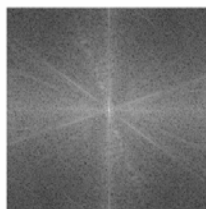
High-boost filter



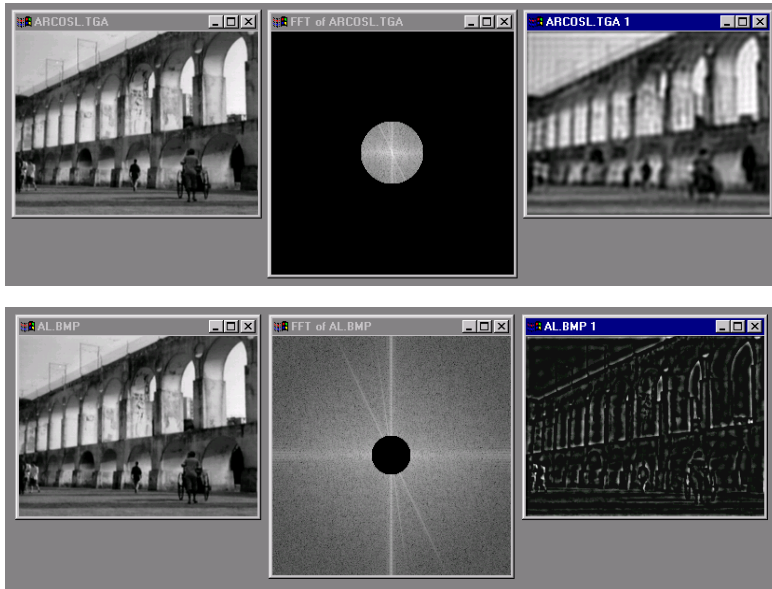
High boosted image



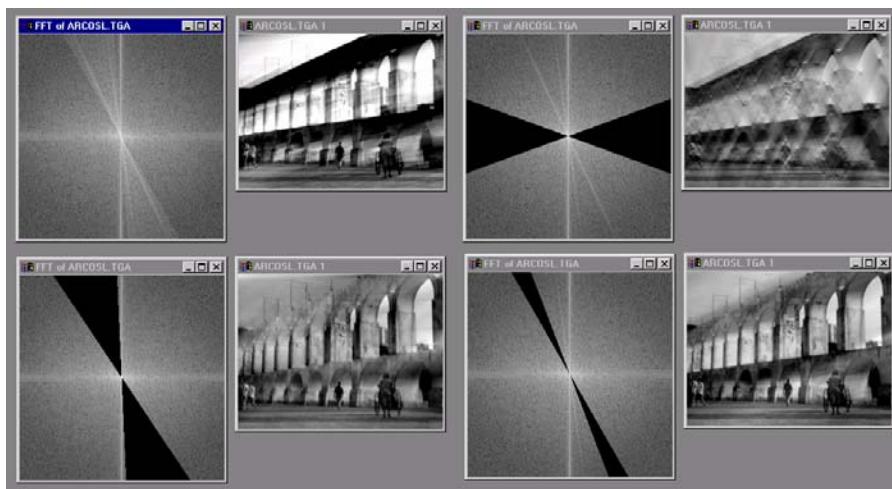
FFT of high boosted image



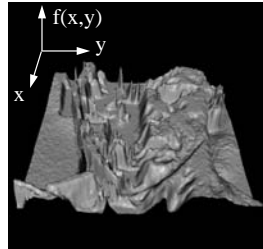
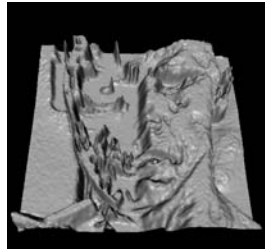
Most information at low frequencies!



## Fun with Fourier Spectra



# Image as a Discrete Function



## Digital Images

The scene is

- **projected** on a 2D plane,
- **sampled** on a regular grid, and each sample is
- **quantized** (rounded to the nearest integer)

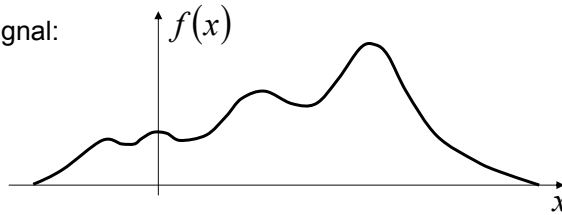
$$f(i, j) = \text{Quantize}\{f(i\Delta, j\Delta)\}$$

Image as a matrix

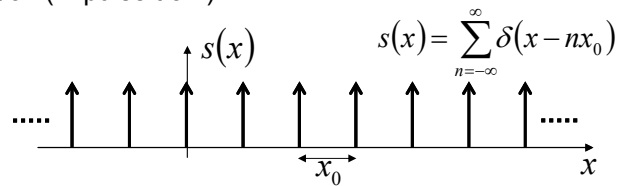
$i \downarrow j \rightarrow$								
	62	79	23	119	120	105	4	0
	10	10	9	62	12	78	34	0
	10	58	197	46	46	0	0	48
	176	135	5	188	191	68	0	49
	2	1	1	29	26	37	0	77
	0	89	144	147	187	102	62	208
	255	252	0	166	123	62	0	31
	166	63	127	17	1	0	99	30

# Sampling Theorem

Continuous signal:



Shah function (Impulse train):



Sampled function:

$$f_s(x) = f(x)s(x) = f(x) \sum_{n=-\infty}^{\infty} \delta(x - nx_0)$$

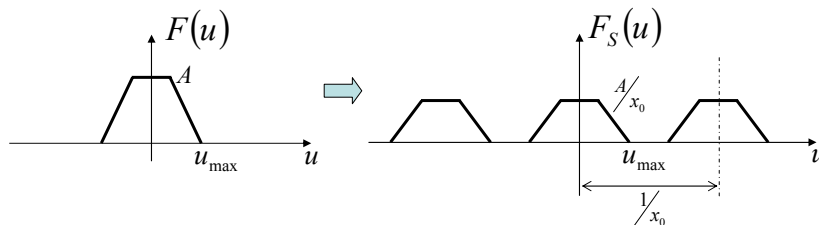
# Sampling Theorem

Sampled function:

$$f_s(x) = f(x)s(x) = f(x) \sum_{n=-\infty}^{\infty} \delta(x - nx_0)$$

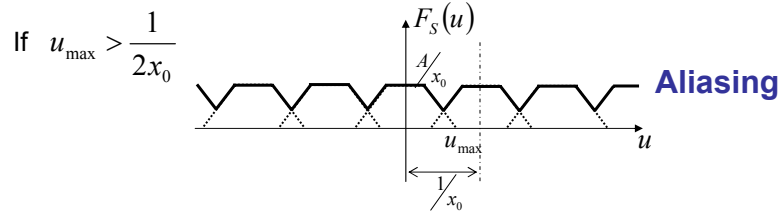
Sampling frequency  $\frac{1}{x_0}$

$$F_s(u) = F(u) * S(u) = F(u) * \frac{1}{x_0} \sum_{n=-\infty}^{\infty} \delta\left(u - \frac{n}{x_0}\right)$$



Only if  $u_{\max} \leq \frac{1}{2x_0}$

# Nyquist Theorem



When can we recover  $F(u)$  from  $F_S(u)$  ?

Only if  $u_{\max} \leq \frac{1}{2x_0}$  (Nyquist Frequency)

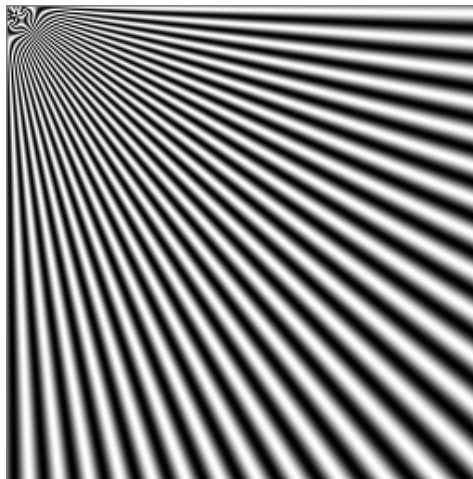
We can use

$$C(u) = \begin{cases} x_0 & |u| < 1/(2x_0) \\ 0 & \text{otherwise} \end{cases}$$

Then  $F(u) = F_S(u)C(u)$  and  $f(x) = \text{IFT}[F(u)]$

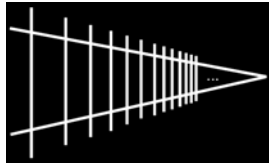
Sampling frequency must be greater than  $2u_{\max}$

# Aliasing



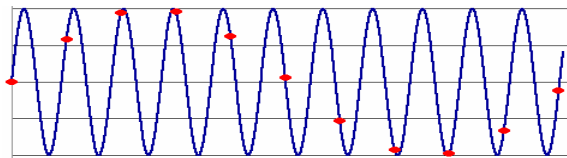


## Alias: n., an assumed name



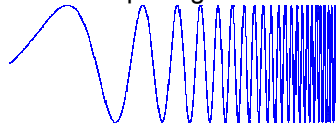
Picket fence receding into the distance will produce aliasing...

WHY?

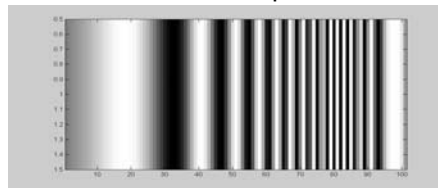


Not enough samples

Input signal:



Matlab output:



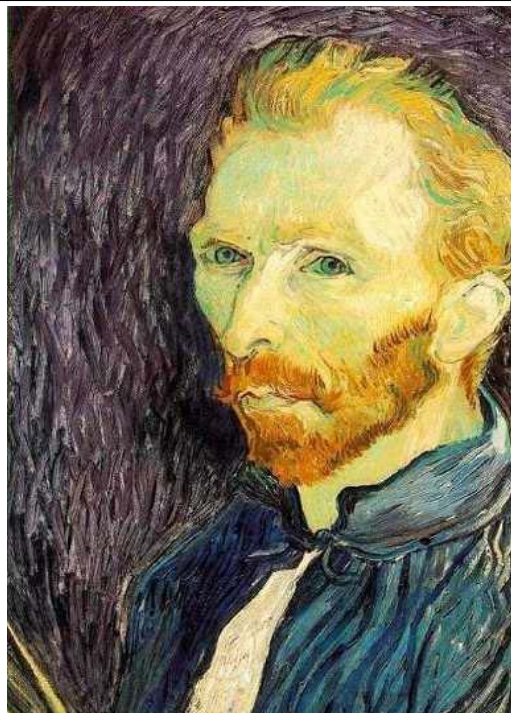
`x = 0:.05:5; imagesc(sin((2.^x).*x))`

Alias!

## Image Scaling

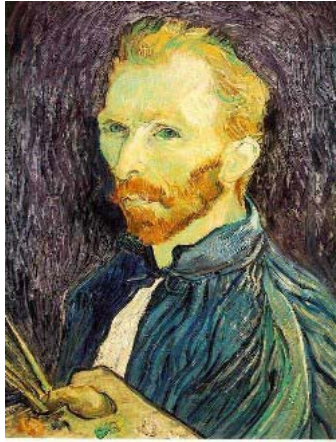
This image is too big to fit on the screen. How can we reduce it?

How to generate a half-sized version?



## Image Sub-Sampling

---



1/4

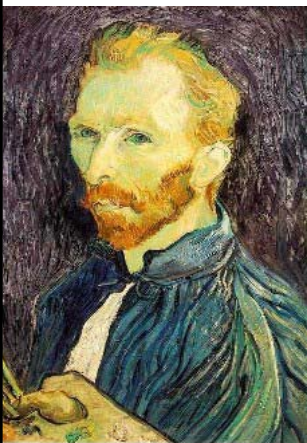


1/8

Throw away every other row and column to create a 1/2 size image  
- called *image sub-sampling*

## Image Sub-Sampling

---



1/2

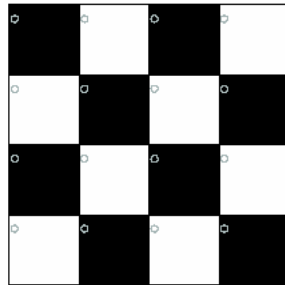
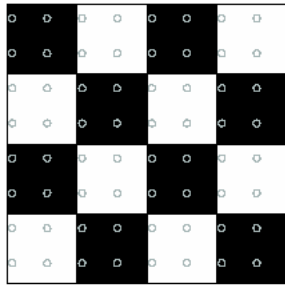


1/4 (2x zoom)

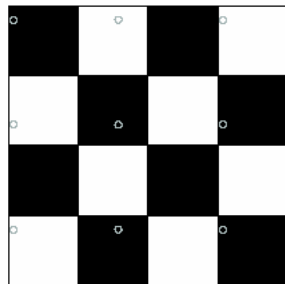
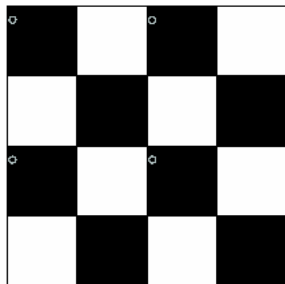


1/8 (4x zoom)

## Good and Bad Sampling



Good sampling:  
• Sample often or,  
• Sample wisely

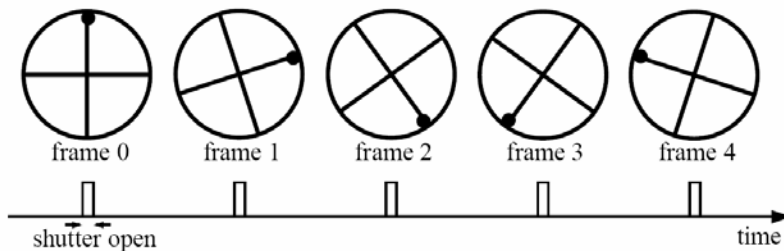


Bad sampling:  
• see aliasing in action!

## Really bad in video

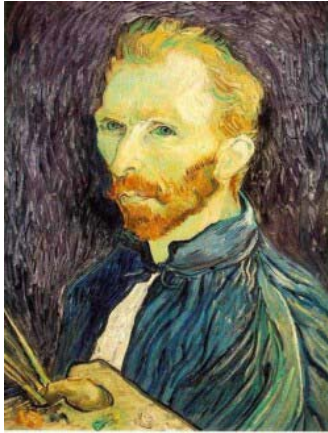
Imagine a spoked wheel moving to the right (rotating clockwise).  
Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time = 1/30 sec. for video, 1/24 sec. for film):



Without dot, wheel appears to be rotating slowly backwards!  
(counterclockwise)

## Sub-Sampling with Gaussian Pre-Filtering



Gaussian 1/2



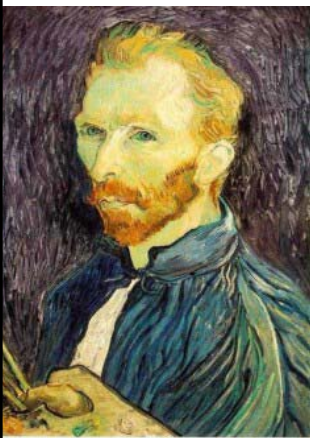
G 1/4



G 1/8

- Solution: filter the image, *then* subsample
  - Filter size should double for each  $\frac{1}{2}$  size reduction. Why?

## Sub-Sampling with Gaussian Pre-Filtering



Gaussian 1/2

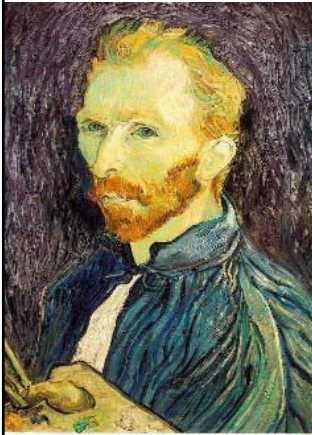


G 1/4



G 1/8

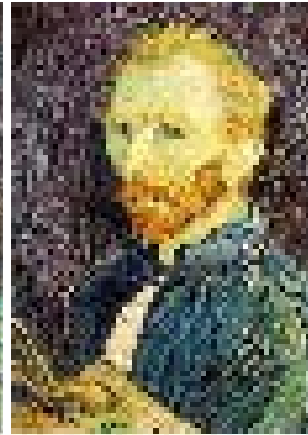
## Compare with...



1/2



1/4 (2x zoom)



1/8 (4x zoom)

## Aliasing

A screenshot of a document viewer window titled 'gv: kpathsea'. The window displays a document with a central toolbar and two side panels. The left panel shows the document's table of contents, with 'Chapter 1: Introduction' and '1 Introduction' selected. The right panel shows the content of the selected section. The text in the right panel is blurry and pixelated, illustrating the effect of aliasing. The toolbar includes buttons for 'File', 'State', 'Page', 'Portrait', 'Fixed Size', 'Open', 'Print All', 'Print Marked', 'Save All', 'Save Marked', 'Redisplay', and a list of page numbers (1-11). The document content includes the following text:

Chapter 1: Introduction

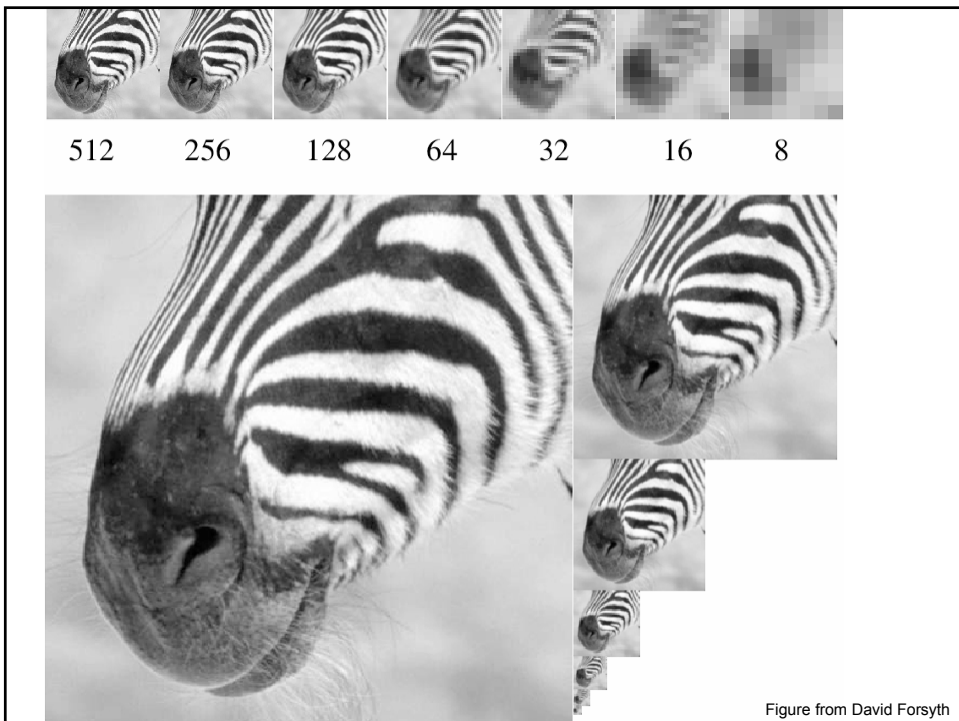
### 1 Introduction

This manual corresponds to version 3.2 of the library's fundamental purpose is to provide a user, similar to what shells do when looking up the following software, all of which we maintain:

- Dvilk (see the 'dvilk' man page)
- Dvipk (see section "Introduction" in Dvilk)
- GNU font utilities (see section "Introduction" in Web2c)
- Web2c (see section "Introduction" in Web2c)
- Xdvik (see the 'xdvik' man page)

Other software that we do not maintain also includes:





## Next Class

---

- Image Processing and Filtering (continued) –  
Edge Detection
- Horn, Chapter 6