

# Computer Vision

Spring 2006 15-385,-685

Instructor: S. Narasimhan

Wean 5403


T-R 3:00pm – 4:20pm

## Image Resampling and Pyramids


### Lecture #8

## RECAP - Sampling Theorem

Continuous signal:



Shah function (Impulse train):



Sampled function:

$$f_s(x) = f(x)s(x) = f(x) \sum_{n=-\infty}^{\infty} \delta(x - nx_0)$$


## RECAP - Sampling Theorem

Sampled function:

$$f_s(x) = f(x)s(x) = f(x) \sum_{n=-\infty}^{\infty} \delta(x - nx_0)$$


Sampling frequency  $\frac{1}{x_0}$

$$F_s(u) = F(u) * S(u) = F(u) * \frac{1}{x_0} \sum_{n=-\infty}^{\infty} \delta\left(u - \frac{n}{x_0}\right)$$



Only if  $u_{\max} \leq \frac{1}{2x_0}$

## RECAP - Nyquist Theorem



When can we recover  $F(u)$  from  $F_S(u)$  ?

Only if  $u_{\max} \leq \frac{1}{2x_0}$  (Nyquist Frequency)


We can use

$$C(u) = \begin{cases} x_0 & |u| < 1/2x_0 \\ 0 & \text{otherwise} \end{cases}$$

Then  $F(u) = F_S(u)C(u)$  and  $f(x) = \text{IFT}[F(u)]$

Sampling frequency must be greater than  $2u_{\max}$


## Aliasing




## Image Scaling

This image is too big to fit on the screen. How can we reduce it?

How to generate a half-sized version?



## Image Sub-Sampling




1/4



1/8

Throw away every other row and column to create a 1/2 size image  
- called *image sub-sampling*

## Image Sub-Sampling





1/2

1/4 (2x zoom)



1/8 (4x zoom)

## Good and Bad Sampling



Good sampling:


- Sample often or,
- Sample wisely



Bad sampling:

- see aliasing in action!

## Sub-Sampling with Gaussian Pre-Filtering



Gaussian 1/2



G 1/4



G 1/8

- Solution: filter the image, *then* subsample
  - Filter size should double for each  $\frac{1}{2}$  size reduction. Why?

## Sub-Sampling with Gaussian Pre-Filtering



Gaussian 1/2




G 1/4



G 1/8


## Compare with...



1/2




1/4 (2x zoom)



1/8 (4x zoom)

## Image Resampling

- What about arbitrary scale reduction?
- How can we increase the size of the image?



- Recall how a digital image is formed


$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale



# Image Resampling


- So what to do if we don't know  $f$ 
  - Answer: guess an approximation  $\tilde{f}$
  - Can be done in a principled way: filtering




- Image reconstruction
  - Convert  $F$  to a continuous function  
 $f_F(x) = F(\frac{x}{d})$  when  $\frac{x}{d}$  is an integer, 0 otherwise
  - Reconstruct by cross-correlation:  
$$\tilde{f} = h \otimes f_F$$

# Resampling filters

- What does the 2D version of this hat function look like?



performs  
linear interpolation



performs  
bilinear interpolation


- Better filters give better resampled images
  - Bicubic is common choice



## Bilinear interpolation

---



- A common method for resampling images



$$\begin{aligned} F(x, y) = & (1-a)(1-b) F(i, j) \\ & + a(1-b) F(i+1, j) \\ & + ab F(i+1, j+1) \\ & + (1-a)b F(i, j+1) \end{aligned}$$

## Image Rotation

---



## Multi-Resolution Image Representation

---

- Fourier domain tells us “what” (frequencies, sharpness, texture properties), but not “where”.
- Spatial domain tells us “where” (pixel location) but not “what”.
- We want a image representation that gives a local description of image “events” – what is happening where.
- Naturally, think about representing images across varying scales.

## Multi-resolution Image Pyramids

---


Low resolution



High resolution




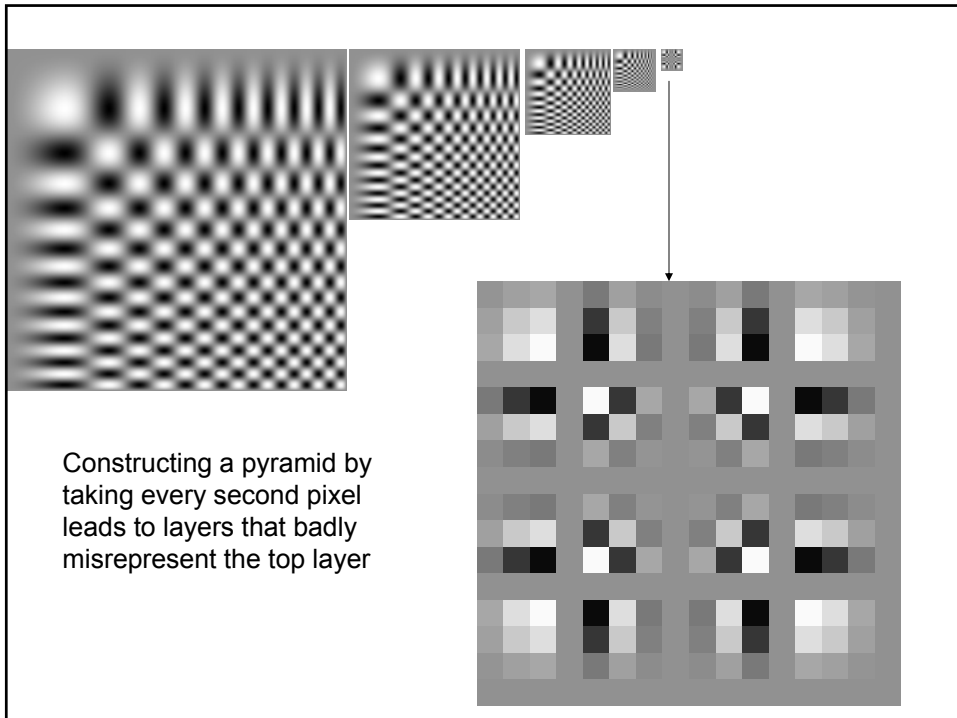
## Space Required for Pyramids



$$N^2 + \frac{1}{4}N^2 + \frac{1}{16}N^2 + \dots = 1\frac{1}{3}N^2$$

## Pyramid construction





Even worse for synthetic images

---



## Decimation




\*

1	4	6	4	1
4	16	24	16	4
6	24	36	24	6
4	16	24	16	4
1	4	6	4	1

Filter





Subsample



## Expansion



Expand  
Image



Interpolation



## Interpolation Results

Original Image



Nearest Neighbor




Bilinear Interpolation




## The Gaussian Pyramid


- Smooth with Gaussians because
  - a Gaussian\*Gaussian=another Gaussian
- Synthesis
  - smooth and downsample
- Gaussians are low pass filters, so repetition is redundant
- Kernel width doubles with each level



Level 1



Level 2

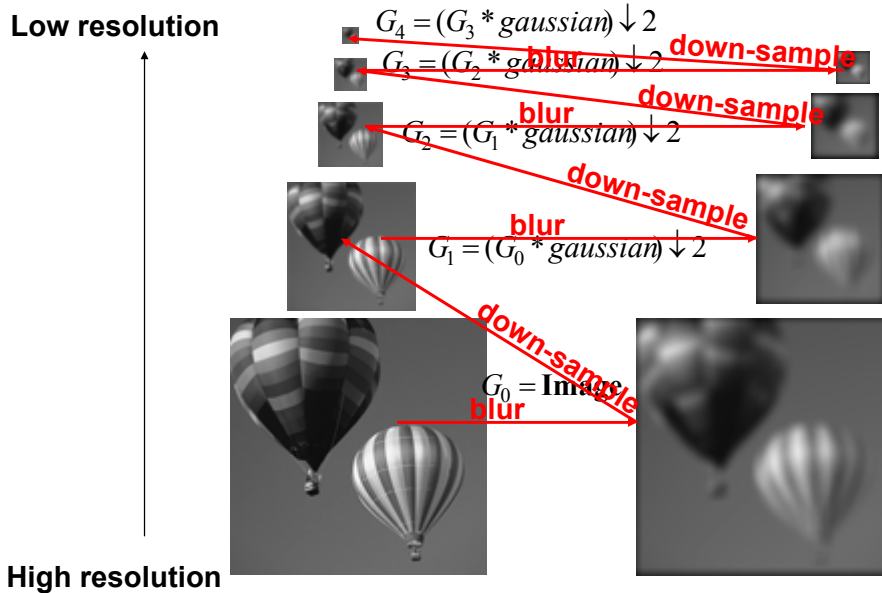


Level 3


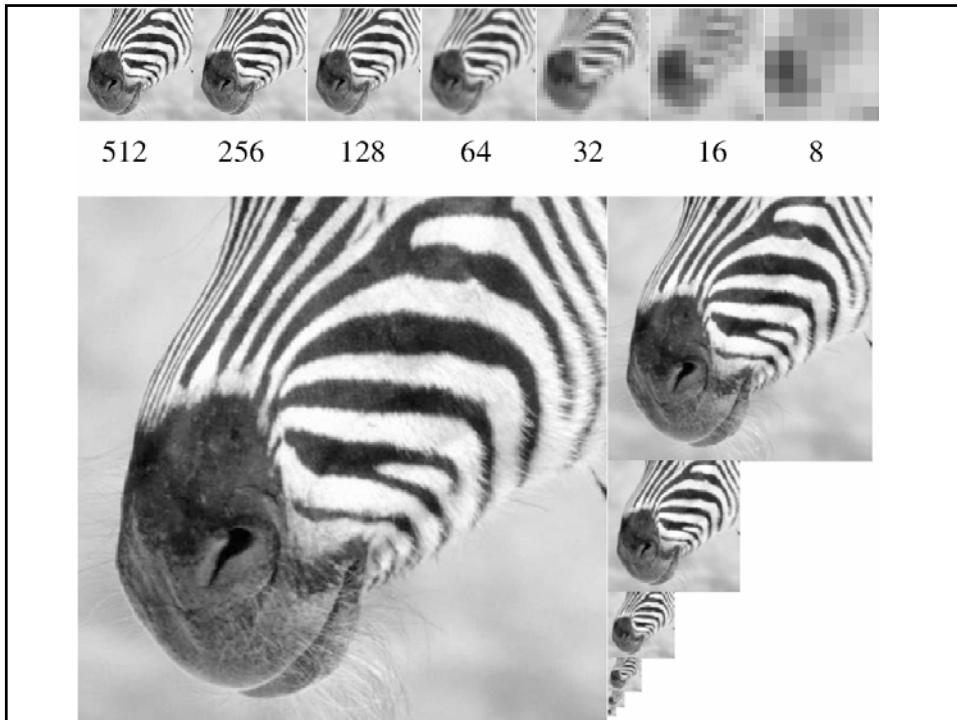
## Smoothing as low-pass filtering

- High frequencies lead to trouble with sampling.
- Suppress high frequencies before sampling !
  - truncate high frequencies in FT
  - or convolve with a low-pass filter
- Common solution: use a Gaussian
  - multiplying FT by Gaussian is equivalent to convolving image with Gaussian.

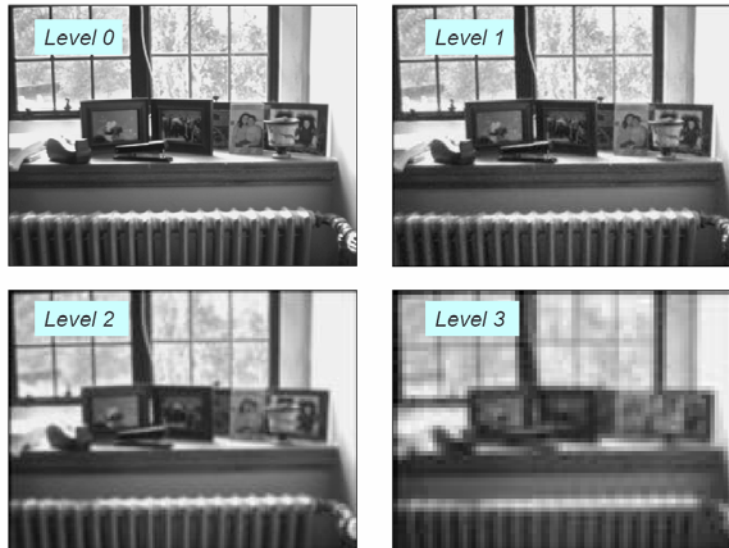
## The Gaussian Pyramid







## Pyramids at Same Resolution



## Difference of Gaussians (DoG)

- Laplacian of Gaussian can be approximated by the difference between two different Gaussians



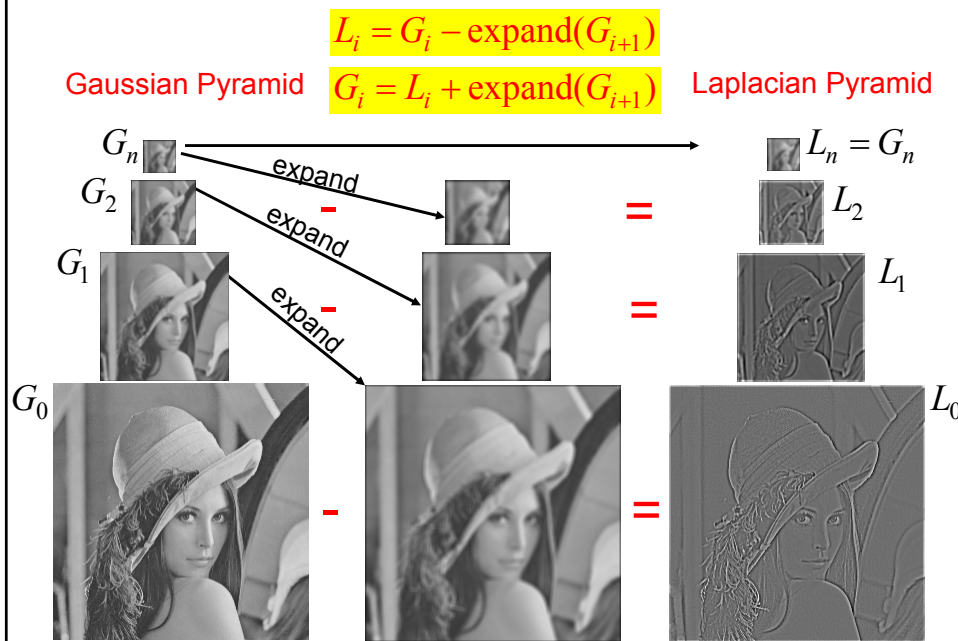



Figure 2-16. The best engineering approximation to  $\nabla^2 G$  (shown by the continuous line), obtained by using the difference of two Gaussians (DOG), occurs when the ratio of the inhibitory to excitatory space constraints is about 1:1.6. The DOG is shown here dotted. The two profiles are very similar. (Reprinted by permission from D. Marr and E. Hildreth, "Theory of edge detection," *Proc. R. Soc. Lond. B* 204, pp. 301-328.)

## Gaussian – Image filter



## The Laplacian Pyramid






## Applications of Image Pyramids

- Coarse-to-Fine strategies for computational efficiency.
- Search for correspondence
  - look at coarse scales, then refine with finer scales
- Edge tracking
  - a “good” edge at a fine scale has parents at a coarser scale
- Control of detail and computational cost in matching
  - e.g. finding stripes
  - very important in texture representation
- Image Blending and Mosaicing
- Data compression (laplacian pyramid)

## Edge Detection using Pyramids

Coarse-to-fine strategy:

- Do edge detection at higher level.
- Consider edges of finer scales only near the edges of higher scales.




## Fast Template Matching

*Template*



*Search Region*



- For an  $m \times n$  image...
- For a  $p \times q$  template...
- The complexity of the 2D pattern recognition task is  $O(mnpq)$  ☹
- This gets even worse for a family of templates (*e.g.*, to address scale and/or rotational effects)

# Fast Template Matching

*Template*



*Search Region*

Original Image





## Multi-resolution correlation

- Multi-resolution template matching
  - reduce resolution of both template and image by creating an **image pyramid**
  - match small template against small image
  - identify locations of strong matches
  - expand the image and template, and match higher resolution template selectively to higher resolution image
  - iterate on higher and higher resolution images
- Issue:
  - how to choose detection thresholds at each level
    - too low will lead to too much cost
    - too high will miss match



## Level 3 Search

- At the lowest pyramid level, we search the entire image with the correlation template



## Level 2 Search



- Subsequent searches are constrained to a neighborhood of only several pixels in the x and y directions







## Level 1 Search

- Subsequent searches are constrained to a neighborhood of only several pixels in the x and y directions



## Level 0 Search

- In the end, the total time (in Matlab) was reduced from  $\approx 31$  seconds to  $\approx 0.5$  seconds while obtaining the same template match




# Image Blending and Mosaicing


- Given two images  $A$  and  $B$
- Construct Laplacian Pyramid  $L_a$  and  $L_b$
- Create a third Laplacian Pyramid  $L_c$  where for each level  $l$

$$L_c(i, j) = \begin{cases} L_a(i, j) & \text{if } i < \text{width} / 2 \\ (L_a(i, j) + L_b(i, j)) / 2 & \text{if } i = \text{width} / 2 \\ L_b(i, j) & \text{if } i > \text{width} / 2 \end{cases}$$

- Sum all levels  $L_c$  in to get the blended image



## Blending Apples and Oranges



## Pyramid blending of Regions


- Given two images  $A$  and  $B$ , and a mask  $M$
- Construct Laplacian Pyramids  $L_a$  and  $L_b$
- Construct a Gaussian Pyramid  $G_m$
- Create a third Laplacian Pyramid  $L_c$  where for each level  $l$

$$L_c(i, j) = G_m(i, j)L_a(i, j) + (1 - G_m(i, j))L_b(i, j)$$

- Sum all levels  $L_c$  in to get the blended image

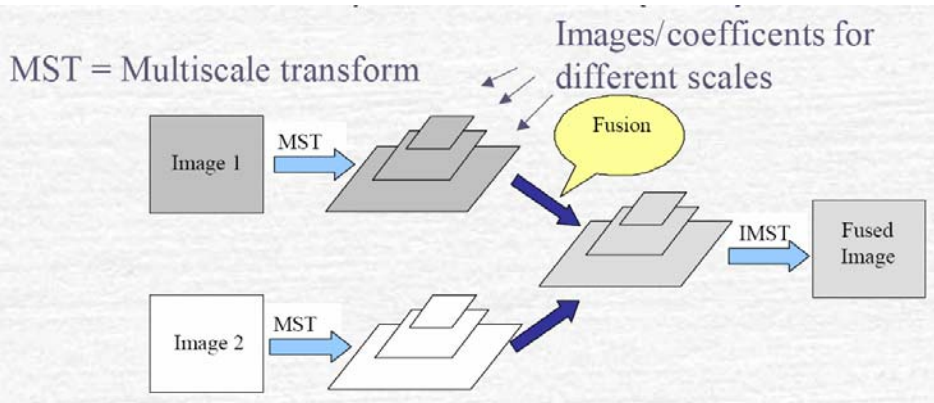


## Horror Photo




© prof. dmartin

## Image Fusion



Multi-scale Transform (MST) = Obtain Pyramid from Image  
Inverse Multi-scale Transform (IMST) = Obtain Image from Pyramid

## Multi-Sensor Fusion



# Image Compression

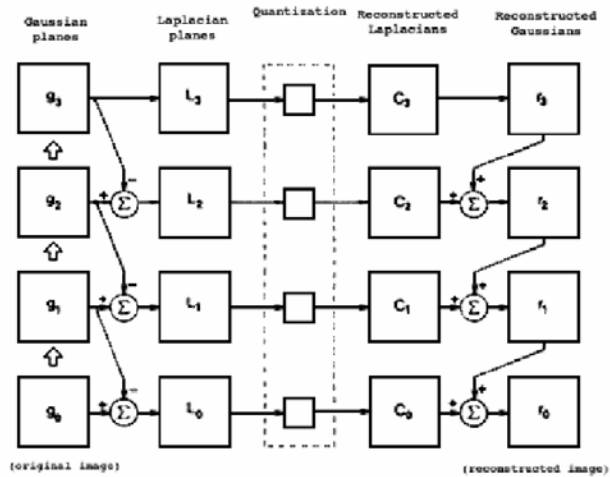


Fig. 10. A summary of the steps in Laplacian pyramid coding and decoding. First, the original image  $g_0$  (lower left) is used to generate Gaussian pyramid levels  $g_1, g_2, \dots$  through repeated local averaging. Levels of the Laplacian pyramid  $L_0, L_1, \dots$  are then computed as the differences between adjacent Gaussian levels. Laplacian pyramid elements are quantized to yield the Laplacian pyramid code  $C_0, C_1, C_2, \dots$ . Finally, a reconstructed image  $r_0$  is generated by summing levels of the code pyramid.

## Next Class

- Hough Transform